


CAS CS 552

Intro to Operating Systems


Richard West

Course Overview and Introduction



Course Summary 1/2

- This course covers the fundamental concepts of operating systems. Topics including OS structure, processes/thread management, synchronization, deadlocks, file systems, I/O and memory management will be discussed
- **A good understanding of C or C++ is required.** A prior understanding of assembly programming will be useful, especially IA32 assembly, but this is not prerequisite knowledge




Course Summary 2/2

- Advanced topics including those based on distributed computing will be discussed, time permitting. Additionally, case studies based on a real-world operating system (e.g., Linux) will be covered, where appropriate, throughout the course
- You will be required to tackle projects that involve kernel-level programming of existing systems such as Linux, and possibly writing your own OS abstractions/features on bare-bones hardware. In the latter case, we will use PC emulation tools such as Bochs, QEMU and VMWare player, time permitting.
 - Our own OS called “Quest” may also be used
- Socket programming and abstractions such as remote procedure calls may be required for some of the projects
- Prerequisites: CS210 (or the consent of the instructor)
 - CS350 is helpful



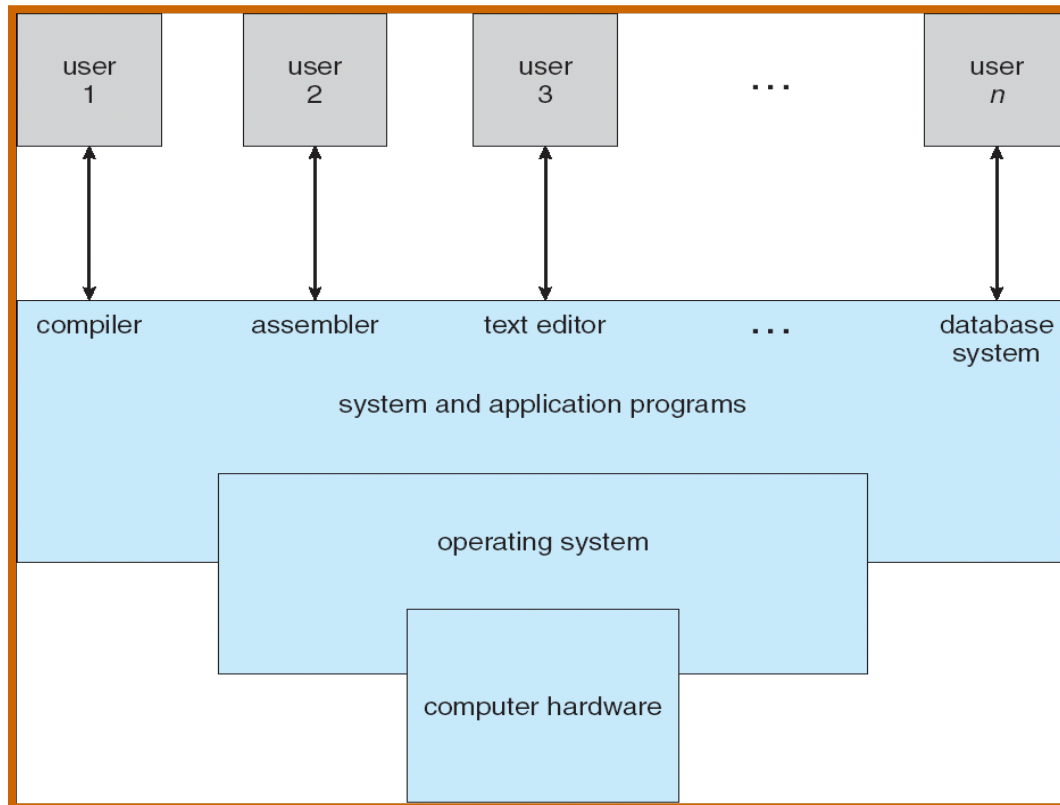
Basic OS Concepts

- Operating system: program between user-level and hardware
 - Allows users to execute progs
 - Aims to provide protection, reliability, efficiency, predictability and convenience
 - “transparency” is a key issue – devil is in the details



Computer Systems

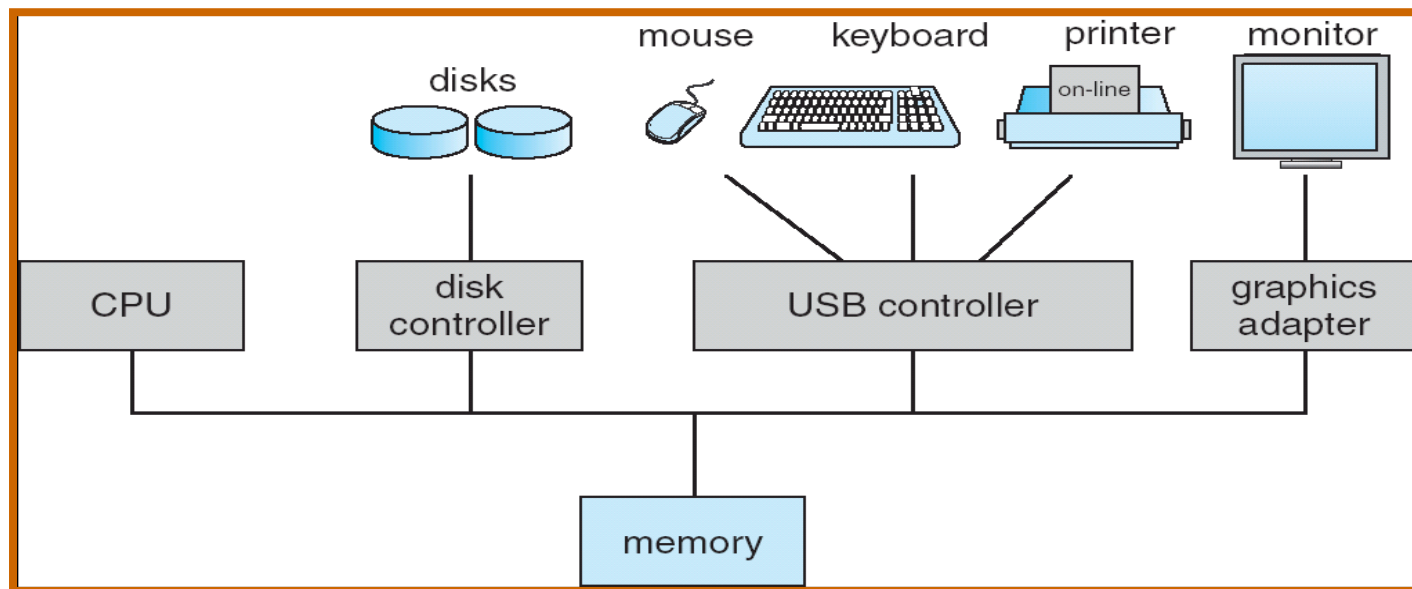
- Computer system: hardware, OS, system/app programs, users



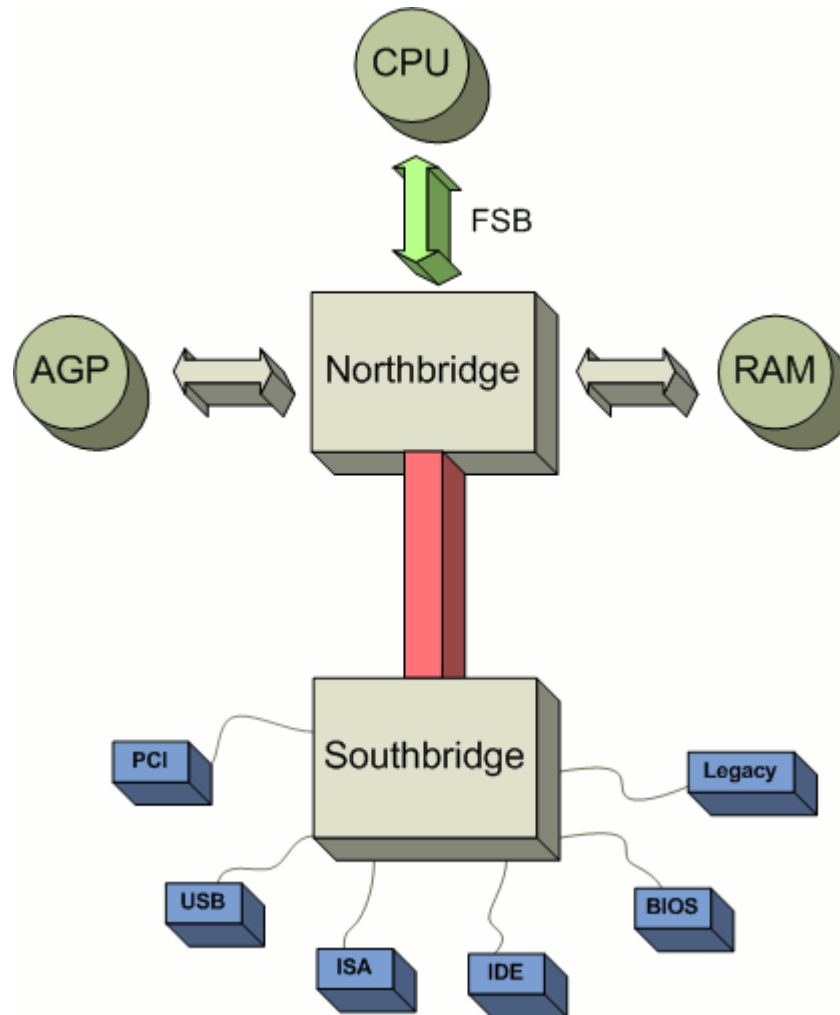
Computer System Organization


Computer-system operation

- One or more CPUs, device controllers connect through common bus providing access to shared memory
- Concurrent execution of CPUs and devices competing for memory cycles



Example – PC Chipset Architecture






Role of OS 1/2

- Controls and coordinates use of hardware amongst multiple progs for various users
- Provides environment for programs to execute
- Acts as a resource allocator
 - Allocates processors, memory, I/O devices, file storage space etc to progs/users as necessary
 - (Typically) must manage competing requests for resources efficiently and fairly
 - Maybe predictably too (more on this later)




Role of OS 2/2

- OSes can be viewed as control programs
 - Control execution of user progs to prevent “improper” use of computer’s resources
 - Provide convenience
 - Make it easier for users to access common resources without awareness of hardware specifics




Multiprogrammed Systems

- Multiprogramming – use of CPU by more than one job (task/process/thread)
 - e.g., while one job is blocked waiting for an I/O device another may use the CPU
 - Aim to keep CPU busy – efficiency issue
 - Multiple jobs are resident in main memory at any given time and CPU(s) is(are) multiplexed between them



Time Sharing Systems

- Timesharing (or multitasking) implies many users simultaneously share computer
 - Like multiprogramming
 - CPU switches between multiple jobs
 - Job switches are frequent, allowing users to interact with program during execution
- User-computer interaction –
 - E.g. input via keyboard, output via display device
 - Typically requires a file system to store data + code
 - Files store a collection of related data e.g., text, programs etc
- Interactive systems require short response times to user requests (typically a few seconds at most)




Personal Computer Systems

- As computer costs decreased, became feasible to have a machine for one user again i.e., personal computers
- PC OSES aim to provide convenience in preference to resource efficiency (or high resource utilization)
 - Possible due to reduced costs of e.g., memory and CPUs
 - Moore's Law
 - PC OSES include Windows (NT,95,98,Me,Vista,7-10), Mac OS X, Linux



- ## Moore's Law – The number of transistors on integrated circuit chips (1971-2018)

This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are linked to Moore's law.






Example Past OSes

- 1951 – LEO 1 (Lyon's UK) for EDSAC
- 1955 – General Motors OS for IBM 701
- 1961 – Burroughs MCP
- 1962 – Atlas Supervisor (for Atlas Supercomputer)
 - One-level store (a.k.a. paged virtual memory)
- 1964-66 – OS/360 IBM, TOPS-10 DEC, Berkeley Timesharing System
- 1965 – T.H.E. Dijkstra et al
- 1965-1970 Multics (MIT)
 - Bell Labs – UNIX
- Circa 1970 UNIX for PDP-11 minicomputer
- (More history later...)
- Today, UNIX-like features in PC OSes




Parallel Systems (1)


- Multiprocessor systems have more than one CPU
- Modern chip multi-processors (CMPs) have multiple cores in the same chip package
- Parallelism goal: to provide computational speedup (of program execution)
- Multiple processors can share prog data, stored on same disks and kept in physically shared memory, rather than spread across multiple machines
- But what about shared caches? NUMA effects? Memory bus bandwidth contention and so forth?






Non-Uniform Memory Access





Parallel Systems (2)

- Provide fault-tolerance/reliability – failure of one CPU only slows down execution rather than resulting in system failure
- Symmetric multiprocessing (SMP) – each processor runs an identical copy of OS
 - CPUs communicate as necessary to coordinate resource management
 - Most modern OSes (Linux, Windows, OS X, Solaris) are SMP-compliant these days
- Asymmetric multiprocessing – each processor is assigned a specific task
 - Master processor controls system and slave processors to perform various tasks/services



Distributed Systems

- Loosely coupled processors on different machines communicate across networks and collectively perform a distributed computation
 - No physically shared memory or common clock
- Enable sharing of remote resources
- Provide computational speedup
- Reliability through replication/redundancy



Real-Time Systems

- Timeliness constraints on operation of a processor or execution of a job/task
 - Program correctness depends not only on the result of execution but when response is generated
 - Hard real-time systems – missing any single time constraint (deadline) can result in failure of the application/system e.g., flight navigation/control, factory automation, nuclear power plant
 - Soft real-time systems – can tolerate certain fraction of deadline misses at cost of reduction in perceived quality/utility e.g., multimedia systems



Other Systems?

- Embedded
 - Often real-time, typically supporting only one application or dedicated function (not general purpose)
 - Found in avionics, automotive and manufacturing environments, medical equipment, and now PDAs, mobile phones...
- Peer-to-peer (P2P)
 - No traditional backend server for a number of clients
 - Each node in distributed system considered a “peer”
 - Popular for information storage, lookup and retrieval
 - e.g., Gnutella, Kazaa, Chord, CAN, Pastry, Bittorrent
 - Made popular as a result of file sharing (especially mp3 audio files)
- Mobile Oses: Android, iOS, WebOS, Symbian, Windows Phone 7 to 8.1, RIM, Maemo, ...