

CAS CS 552 Intro to Operating Systems


Richard West

Processes

Process Concept

- Process – a program in a state of execution
 - A program is a passive entity. It may consist of code and data in files, which are stored on disk and brought to main memory when the program is executed
 - A process is an active entity and includes:
 - a program counter value
 - stack
 - data segment
 - text segment
 - heap


Process in Memory



Process State

- As a process executes, it changes *state*
- Possible states:
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting/blocked**: The process is waiting for some event to occur (e.g., I/O completion, signal from OS)
 - **ready**: The process is runnable but waiting to be assigned to a CPU
 - **terminated**: The process has finished execution


Diagram of Process State



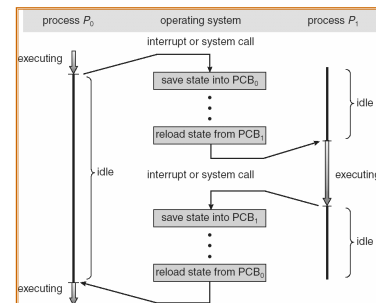
Process Control Block (PCB)

- OS maintains a process/task control block (PCB) for each process in system
- Information associated with each process includes:
 - Process (run) state
 - Program counter value
 - CPU registers
 - CPU scheduling information – priority, size of time-slice, pointers to scheduling queues etc...
 - Memory-management information – base/limit values for address space, page table info, list of virtual memory regions that make up address space...
 - Accounting information – how long process has been running, time spent at user/kernel level etc...
 - File management info – list of open files in use by process...
 - I/O status information – list of devices allocated to a process...

Process Control Block (PCB)



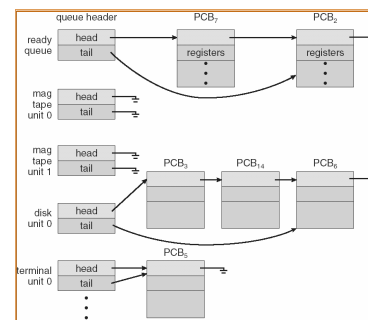
CPU Switch From Process to Process




Process Scheduling Queues

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - Processes are dispatched from the ready queue and assigned to the CPU
- **Device queues** – set of processes waiting for an I/O device
 - Typically one device queue for each device
- Processes migrate among the various queues during their lifetimes

Ready Queue And Various I/O Device Queues



Representation of Process Scheduling



- During the execution of a process, P , several events may occur that cause P to be moved to a wait queue, or the "ready queue"

Schedulers

- **Long-term scheduler** – selects which processes should be brought into the ready queue
 - Controls degree of multiprogramming by selecting processes from disk to be loaded into main memory
- **Short-term scheduler** (or CPU scheduler) – selects a process for execution from the ready queue and allocates CPU

CPU- versus I/O-bound Processes

- Processes can be described as either:
 - I/O-bound processes**
 - spend more time doing I/O than computations, many short CPU bursts
 - CPU-bound processes**
 - spend more time using CPU cycles rather than waiting on I/O
- Long term scheduling can be used to ensure "good mix" of CPU- and I/O-bound processes are present in main memory

Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
- Context-switch time is overhead; the system does no useful work while switching
- Time dependent on hardware support:
 - Memory speed, machine state (#registers)...
 - Nowadays, in the nano-/microsecond range

Operations on Processes

- Process Creation
 - Parent processes may create/fork child processes
 - Children may inherit/share system resources from/with parent process (e.g., some memory, open files...)
- When a parent creates a child process:
 - Parent and child may execute concurrently
 - Parent can wait until some/all children have terminated
 - e.g., fork() versus vfork() on UNIX systems

UNIX Process Creation

- fork() creates "copy" of parent's address space for child...or does it?
 - What about "copy on write" (COW)?
 - When is COW most useful?
- An exec() call (typically execve()) replaces address space of caller with new program image
- Extra: vfork() suspends parent while child uses parent's page table mappings. Parent resumes either when child calls exec() or _exit()...child cannot call exit()!

Process Termination

- exit() routine terminates process, releasing its resources back to system
 - On UNIX systems, _exit() is syscall that releases memory allocated to process
 - May need to keep enough state for a "dead" process in case parent wishes to find exit status via a wait/waitpid() call
 - Can lead to "zombie" processes

Additional Information

- The following slides are taken, for the most part, from Silberschatz et al...


Process Creation

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate

Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace the process' memory space with a new program

Process Creation

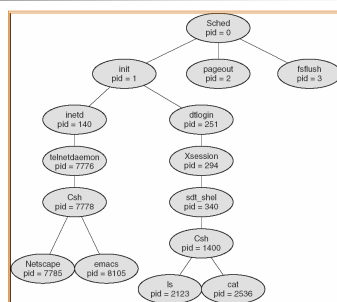


C Program Forking Separate Process

```

int main()
{
    Pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
        complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
    
```

A tree of processes on a typical Solaris



Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**)
 - Termination status of child sent to parent (via **wait**)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort**)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - Some operating systems do not allow child to continue if its parent terminates
 - All children terminated - *cascading termination*

Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - *unbounded-buffer* places no practical limit on the size of the buffer
 - *bounded-buffer* assumes that there is a fixed buffer size

Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```
- Solution is correct, but can only use BUFFER_SIZE-1 elements

Bounded-Buffer – Insert() Method

```
while (true) {
    /* Produce an item */
    while (((in = (in + 1) % BUFFER_SIZE)
count) == out)
        ; /* do nothing -- no free buffers */
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
}
```

Bounded Buffer – Remove() Method

```
while (true) {
    while (in == out)
        ; // do nothing -- nothing to
consume

    // remove an item from the buffer
    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    return item;
}
```

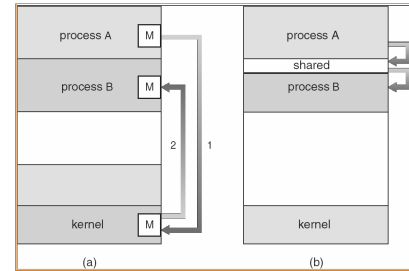
Interprocess Communication (IPC)

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send(*message*)** – message size fixed or variable
 - **receive(*message*)**
- If *P* and *Q* wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

Communications Models



Direct Communication

- Processes must name each other explicitly:
 - send**(P , *message*) – send a message to process P
 - receive**(Q , *message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - send**(A , *message*) – send a message to mailbox A
 - receive**(A , *message*) – receive a message from mailbox A

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue
 - **Non-blocking receive** has the receiver receive a valid message or null

Buffering

- Queue of messages attached to the link; implemented in one of three ways
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits

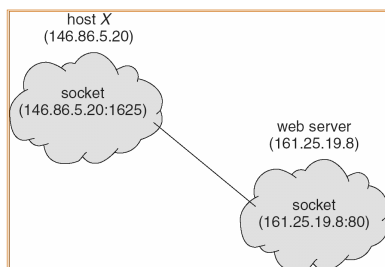
Client-Server Communication

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)

Sockets

- A socket is defined as an *endpoint for communication*
- Concatenation of IP address and port
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets

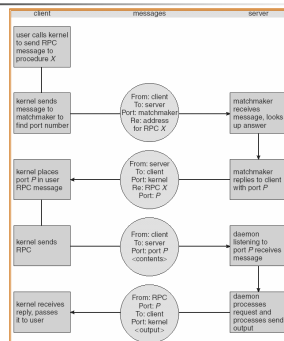
Socket Communication



Remote Procedure Calls


- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
- **Stubs** – client-side proxy for the actual procedure on the server
- The client-side stub locates the server and *marshalls* the parameters
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server

Execution of RPC



Remote Method Invocation

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs.
- RMI allows a Java program on one machine to invoke a method on a remote object.



Marshalling Parameters

