



# CAS CS 552

## Intro to Operating Systems

---

Richard West

Synchronization



# Process Synchronization

---

- Cooperating processes share data (or, more generally, they share resources)
- Concurrent access to shared data may result in data inconsistency (or, more generally, unacceptable interleaved access to resources)



# Example: Producer-Consumer Bounded Buffer 1/2

---

## producer

```
while (true) {  
    produce item;  
    while (counter==n);  
    buffer[in] = item;  
    in = (in+1)%n;  
    counter++;  
}
```

## consumer

```
while (true) {  
    while (counter==0);  
    item = buffer[out];  
    out = (out+1)%n;  
    counter--;  
    consume item;  
}
```



## Example: Producer-Consumer 2/2

---

- Suppose producer & consumer execute concurrently
  - e.g., let counter = 10 and producer/consumer execute counter++ and counter--, respectively
  - After these statements, counter may be 9, 10, or 11, although the correct answer is 10
  - counter = 10 is only guaranteed if producer & consumer execute the instructions that modify counter separately (i.e., atomically!)



# Synchronization Principle

---

- Arbitrary interleaving of the execution of concurrent processes when accessing shared data can lead to data inconsistency
  - What we have is a “data race condition”
- We must ensure that cooperating processes access shared data one at a time using some method that enforces synchronization



# Critical Sections

---


- Critical sections = code segments which access shared data/resources
  - Only one process can execute a critical section at a time
  - Critical sections require mutually exclusive access



# Dealing with the Critical Section Problem

---

- (1) no two processes may be together in a critical section (C.S.)
- (2) no assumptions should be made about the relative speeds of processes
- (3) no process outside its critical section should block (the progress) of other processes
- (4) no process should wait arbitrarily long to enter its critical section



# Two-Process Solution to the C.S. Problem

---

- Strict alternation (method 1)
- Consider two processes,  $P_i$  and  $P_j$

```
while (true) {  
    non-critical section;  
    while (turn!=i);  
    critical section;  
    turn = j;  
}
```

Problem: Does not satisfy rule (3) –  $P_i$  cannot enter its critical section if  $\text{turn} \neq i$ , even though  $P_j$  may not be in its critical section





# Better Solution to C.S. Problem with Two Processes 1/2

---

- (method 2) introduce a boolean variable “flag”
  - `typedef enum {false, true} boolean;`
  - `boolean flag[2];`
  - `enum {0,1} turn;`
- Initially, `flag[0]=flag[1]=false;`
  - turn can be either 0 or 1;



# Better Solution to C.S. Problem with Two Processes 2/2

---

```
while (true) {  
    non-critical section;  
    flag[i]=true;  
    turn=j;  
    while ((flag[j]) && (turn==j));  
    critical section;  
    flag[i]=false;  
}
```

- $i$  = local process ID;  $j$  = remote/other process ID
- This solution satisfies all four conditions for dealing with critical section problem



# Multiple Process Solutions

---

- Bakery algorithm for  $n$  processes
  - Customers assigned ticket numbers
  - Customer with lowest ticket number is serviced next
  - Two customers can have same ticket number, so ties are broken according to some rule:
    - e.g., customer  $C_i$  precedes  $C_j$  if  $ID(C_i) < ID(C_j)$
    - This requires customers to be numerically and uniquely identified



# Bakery Algorithm 1/2

---

- Common data structures:
  - Boolean choosing[n];
  - Int number[n];
  - Initially, these data structures are set to **false** and **0**, respectively
- Notation:
  - Let,  $(a,b) < (c,d)$  if  $(a < c) \parallel ((a == c) \ \&\& \ (b < d))$
  - $\max(a_0, \dots, a_{n-1})$  = largest value in set  $\{a_0, \dots, a_{n-1}\}$
- The following guarantees all four conditions for dealing with critical sections...



# Bakery Algorithm 2/2

---

- Process  $P_i$ :

```
while(true) {  
    choosing[i]=true;  
    number[i]=max(number[0],...,number[n-1])+1;  
    choosing[i]=false;  
    for (j=0;j<n;j++) {  
        while (choosing[j]);  
        while ((number[j]!=0) &&  
                ((number[j],j) < (number[i],i)));  
    }  
    critical section;  
    number[i]=0;  
    non-critical section;  
}
```



# Synchronization Hardware

---

- On uniprocessors, disabling interrupts guarantees safety when a process is executing in a critical section
  - Disabling interrupts is not always practical on a multiprocessor
- Hardware provides atomic instructions
  - Cannot be interrupted but must instead be executed to completion, thereby preventing *interleaved* execution



# Test-and-set

---

- Implemented as an atomic instruction, test-and-set conceptually looks like:

```
atomic boolean test-and-set (boolean *target) {  
    boolean b;  
    b = *target;  
    *target = true;  
    return b;  
}
```




# A test-and-set Mutex

---

```
while (true) {  
    while (test-and-set(&lock));  
    critical section;  
    lock=false;  
    non-critical section;  
}
```

- Initially, lock=false;





# Test-and-set Usage 1/2

---

- Using a shared variable “flag” initially 0, here is another solution to the critical section problem:

enter\_region:

```
    tsl register, flag ; copy flag to register and set  
                        ; flag to 1
```

```
    cmp register, #0 ; was flag 0?
```


```
    jnz enter_region ; if it was non-zero, lock was set  
                        ; so loop
```

```
    ret                ; return to caller
```

leave\_region:

```
    mov flag, #0      ; flag=0
```

```
    ret                ; return to caller
```



## Test-and-set Usage 2/2

---

- Each process has the following pseudo-code:

```
while (true) {  
    enter_region;  
    critical section;  
    leave_region;  
    non-critical section;  
}
```



# Compare-and-Swap (CAS)

---

- e.g., CMPXCHG x86 instruction
- Compare contents of memory location w/ specified value
  - if same, change memory location contents to new value

```
atomic int CAS (int *pval, int old, int new) {  
    int val = *pval;  
    if (val == old)  
        *pval = new;  
    return val;  
}
```



# CAS Usage

---

- Atomic variable increment (fetch-and-add)

```
void incr (int *pval) {  
    do {  
        int tmp = *pval;  
    }  
    while (CAS (pval, tmp, tmp+1) != tmp);  
}
```



# Semaphores

---

- Used for synchronization
- A semaphore “s” is an integer variable accessed via two atomic operations: wait and signal
  - wait (also known as “p” for “proberen” – Dutch for “to test”)
  - signal (also known as “v” for “verhogen” – Dutch for “to increment”)
- Conceptually:
  - `wait(s): while (s<=0); s--;`
  - `signal(s): s++;`



# Semaphores Example

---

- Let “mutex” be a binary semaphore, initialized to 1
- Process,  $P_i$ :

```
while (true) {  
    wait(mutex);  
    critical section;  
    signal(mutex);  
    non-critical section;  
}
```
- All mutual exclusion problems so far require busy waiting
  - While one process is in its critical section another process must loop through its entry code before entering its own critical section



# Spinning versus Blocking Locks

---

- Spinlocks are essentially busy waiting semaphores
  - A process spins in a tight loop waiting for a lock
- Busy waiting (and, hence, spinlocks) are not good for uniprocessor systems – Why?
  - Waste CPU cycles
- Spinlocks *are* useful in multiprocessor systems (and uniprocessor systems with preemptible kernels)
  - No context-switch is required when a process or thread waits for a lock
- Instead of spinning, a process/thread can be blocked so that the CPU can be used by another process/thread



# Blocking Semaphores 1/2

---

- Blocked processes are moved to a wait queue for a given semaphore
  - Blocked processes are moved to the ready queue by a wakeup operation, when a process executes a signal operation on the corresponding semaphore





# Blocking Semaphores 2/2

---

- `wait(s):`  
    `s--;`  
    `if (s < 0) {`  
        `add calling process to wait_queue(s);`  
        `block process;`  
    `}`
- `signal(s):`  
    `s++;`  
    `if (s <= 0) {`  
        `remove calling process from wait_queue(s);`  
        `wakeup process;`  
    `}`



# Semaphore Semantics

---

- If a semaphore  $S$  is negative...
  - Its magnitude indicates the # of processes waiting on  $S$
- If a semaphore  $S$  is positive...
  - Its value indicates the # of instances of the resource guarded by  $S$  that can be acquired concurrently
- If  $S$  has just two values it is a binary semaphore, else it is a counting semaphore



# Deadlocks and Starvation 1/3

---

- Consider the following:

- $P_0$ :

wait(s);

wait(q);

...

signal(s);

signal(q);

- $P_1$ :

wait(q);

wait(s);

...


signal(q);

signal(s);

Can end up with a cycle of waiting processes  $\Rightarrow$  deadlock!

# Deadlocks and Starvation 2/3

- Let:
- $P_i \rightarrow R \Rightarrow$  process  $P_i$  is waiting for a resource guarded by semaphore  $R$
- $R \rightarrow P_i \Rightarrow$  a resource guarded by semaphore  $R$  is allocated to  $P_i$
- A deadlock occurs when there is a cycle as follows:





# Deadlocks and Starvation 3/3

---

- Related to deadlocks is the notion of starvation
  - Starvation occurs when a process cannot get a resource (e.g., CPU, semaphore) and waits indefinitely while other processes may make progress
    - e.g., a LIFO-ordered semaphore wait queue may cause starvation



# Counting Semaphores

---

- Given two binary semaphores, s1 and s2, we can implement a counting semaphore, S, as follows:  

```
binary_semaphore s1=1, s2=0;  
int c; // set to initial value for s
```
- wait(S):  

```
wait(s1);  
c--;  
if (c < 0) {signal(s1); wait(s2);}  
signal(s1);
```
- signal(S):  

```
wait(s1);  
c++;  
if (c <= 0) signal(s2);  
else signal(s1);
```



# Classic Problems

---

- Readers-writers problem:
  - Multiple concurrent processes wish to access a shared data object
    - Some only wish to read – readers
    - Some wish to read & write – writers
- One approach to problem:
  - Allow read access by multiple readers
  - No reader kept waiting unless a writer has already acquired exclusive access to the shared object
- NOTE: writers may starve if object is always being accessed by readers



# Readers-writers Problem

---

```
semaphore: mutex, wrt; //  
    initially, both 1  
int readcount=0;
```

- writer

```
    wait(wrt);  
    update_object;  
    signal(wrt);
```

- reader

```
    wait(mutex);  
    readcount++;  
    if (readcount==1)  
        wait(wrt);  
    signal(mutex);  
    read_object;  
    wait(mutex);  
    readcount--;  
    if (readcount==0)  
        signal(wrt);  
    signal(mutex);
```





# Monitors

---

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time


```
monitor monitor-name {  
    // shared variable declarations  
    procedure P1 (...) { .... }  
    ...  
  
    procedure Pn (...) {.....}  
  
    Initialization code ( ....) { ... }  
    ...  
}  
}
```




# Monitors (continued)

---


- Procedures defined within a monitor can only access those variables declared locally within the monitor, as well as procedure arguments
- Processes waiting to enter the monitor are queued while one process is active
- Monitors support condition variables
  - e.g., `condition_t x,y;`
  - Operations:
    - `x.wait();` // process is suspended until another process invokes `x.signal();`
    - `x.signal();` // resume one suspended process (if any exist) that previously invoked `x.wait()`



# Monitor with Condition Variables



# Dining-Philosophers Problem 1/2



- 5 philosophers, 5 chopsticks and a bowl of rice (data set)



## Dining-Philosophers Problem 2/2

---

- Each philosopher may be in one of 3 possible states
  - **thinking, hungry, eating**
- A **hungry** philosopher tries to acquire one chopstick at a time on his/her left- and right-hand sides
- Only if both chopsticks closest to a philosopher are not being used by another philosopher can they both be picked up
  - At such a time, a philosopher may eat
  - After eating, both chopsticks are placed on the table and the philosopher returns to thinking



# Dining-Philosophers Solution 1/3

---

```
monitor DP {  
    enum {THINKING, HUNGRY, EATING} state[5];  
    condition_t self[5]; //for hungry philosophers waiting to eat  
  
    void pickup(int i) {  
        state[i] = HUNGRY;  
        test(i);  
        if (state[i] != EATING) self[i].wait;  
    }  
  
    void putdown(int i) {  
        state[i] = THINKING;  
        // test left and right neighbors  
        test((i + 4) % 5);  
        test((i + 1) % 5);  
    }  
}
```



# Dining-Philosophers Solution 2/3

---

```
void test(int i) {  
    if ( (state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING;  
        self[i].signal();  
    }  
}  
  
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}  
} // end of monitor
```



# Dining-Philosophers Solution 3/3

---

- Each philosopher  $i$  invokes the operations `pickup()` and `putdown()` in the following sequence:

`dp.pickup(i)`

EAT

`dp.putdown(i)`

- This solution guarantees no two neighbors are eating simultaneously and no deadlocks will occur...but starvation is still possible





# Monitor Implementation Using Semaphores

---

- For each monitor introduce the following variables  
    semaphore\_t mutex; // (initially = 1)  
    semaphore\_t next; // (initially = 0)  
    int next-count = 0; // # processes waiting to go next  
        in monitor
- Each procedure P in monitor will be replaced by  
    wait(mutex);  
        P(); // the actual procedure P  
    if (next-count > 0)  
        signal(next)  
    else  
        signal(mutex);
- Mutual exclusion within a monitor is ensured



# Monitor Implementation (cont.)

---

- For each condition variable  $x$ , we have:

```
semaphore_t x-sem; // (initially = 0)
int x-count = 0;
```

- The operation  $x.wait$  can be implemented as:

```
x-count++;
if (next-count > 0)
    signal(next); // allow a suspended process to
                // resume
else
    signal(mutex); // allow a process to enter
                // monitor
wait(x-sem);
x-count--;
```



## Monitor Implementation (cont.)

---

- The operation **x.signal** can be implemented as:

```
if (x-count > 0) {  
    next-count++;  
    signal(x-sem);  
    wait(next);  
    next-count--;  
}
```



# Synchronization Examples

---

- Solaris
- Windows XP
- Linux
- Pthreads



# Solaris Synchronization

---

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing
- Uses **adaptive mutexes** for efficiency when protecting data from short code segments
- Uses **condition variables** and **readers-writers** locks when longer sections of code need access to data
- Uses **turnstile** to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock



# Windows XP Synchronization

---

- Uses interrupt masks to protect access to global resources on uniprocessor systems
- Uses **spinlocks** on multiprocessor systems
- Also provides **dispatcher objects** which may act as either mutexes and semaphores
- Dispatcher objects may also provide **events**
  - An event acts much like a condition variable



# Linux Synchronization

---

- Linux:
  - disables interrupts to implement short critical sections
- Linux provides:
  - semaphores
  - spin locks




# Pthreads Synchronization

---

- Pthreads API is OS-independent
- It provides:
  - mutex locks
  - condition variables
- Non-portable extensions include:
  - read-write locks
  - spin locks





# Case Study: Synchronization

## (CPU Protection Problem)



# Mutual Exclusion via Semaphores

---

- Ensure only one task/process is in the critical section

- wait(S) to get access to semaphore S

- signal(S) to release S

- Example: producer consumer

```
    Producer()
```

```
    {
```

```
        .
```

```
        .
```

```
        wait(S)
```

```
        critical section /* create data & increment pointer */
```


```
        signal(S)
```

```
        .
```

```
        .
```


```
    }
```

# Priority Inversion



Source of the figure: Chenyang Lu, Washington University Saint Louis

# Unbounded Priority Inversion





# What really happened on Mars?

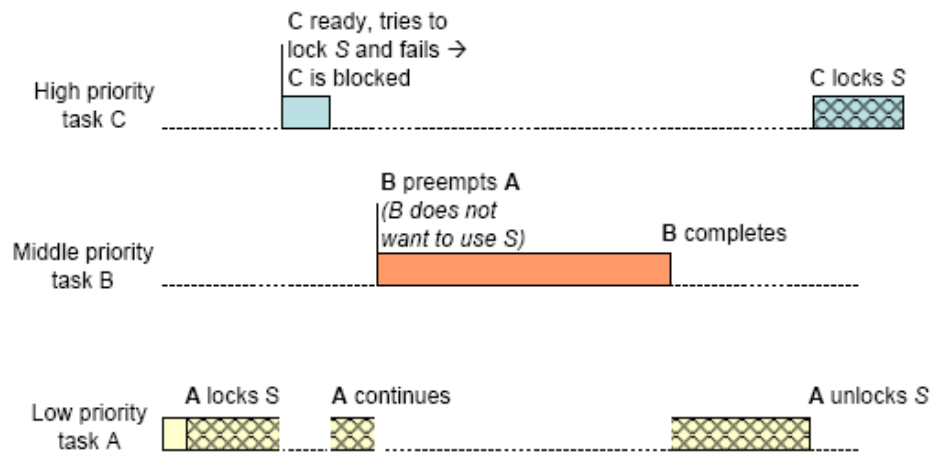
---

## ■ Repeated resets in the Mars Pathfinder

- ❑ The Mars Pathfinder mission was widely proclaimed as "flawless" in the early days after its July 4th, 1997 landing on the Martian surface.... But a few days into the mission, not long after Pathfinder started gathering meteorological data, the spacecraft began experiencing total system resets, each resulting in losses of data. The press reported these failures in terms such as "software glitches" and "the computer was trying to do too many things at once"....
- ❑ For a full story, visit [http://research.microsoft.com/%7Embj/Mars\\_Pathfinder/Mars\\_Pathfinder.html](http://research.microsoft.com/%7Embj/Mars_Pathfinder/Mars_Pathfinder.html)

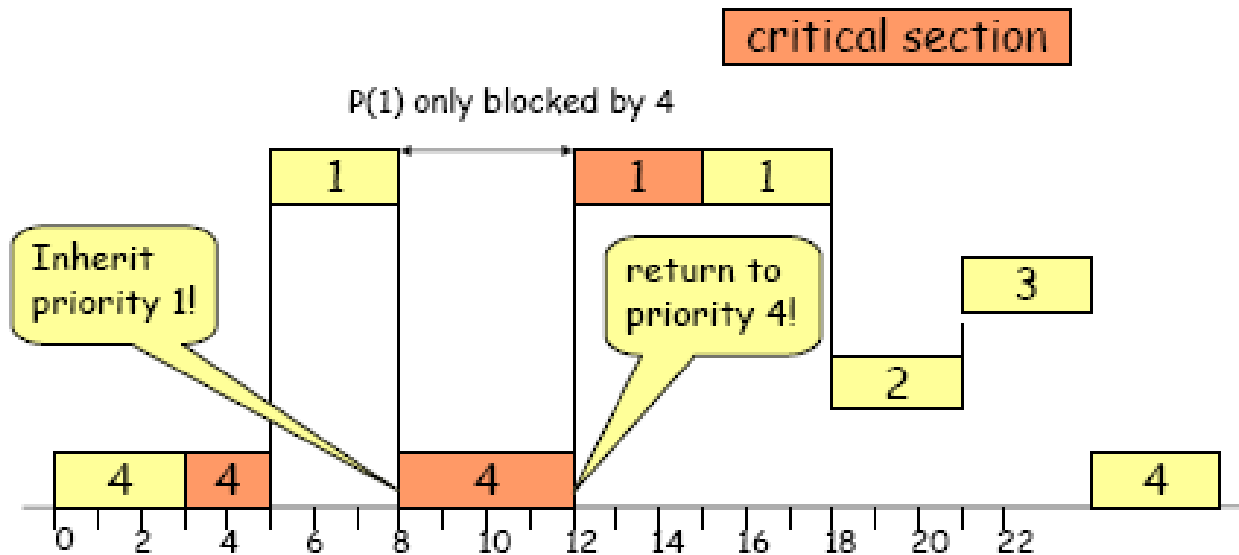
# Pathfinder Incident

- Classical priority inversion problem due to shared system bus!



# Priority Inheritance

- Inherit the priority of the blocked high priority task





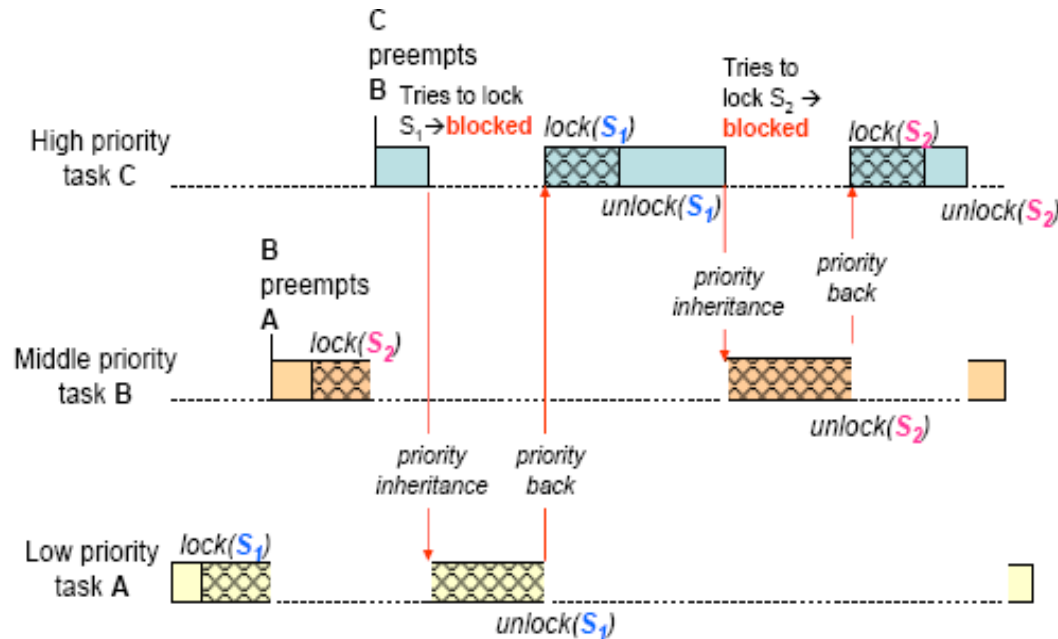
# Priority Inheritance Protocol (PIP)

---

- If  $T_L$  blocks a higher priority task  $T_H$ ,  $\text{priority}(T_L) \leftarrow \text{priority}(T_H)$
- When  $T_L$  releases a semaphore:
  - Return to its normal priority if it doesn't block any task
  - Otherwise, set  $\text{priority}(T_L) \leftarrow$  highest priority of the tasks blocking on a semaphore held by  $T_L$
- Transitive
  - $T_1$  blocked by  $T_2$ :  $\text{priority}(T_2) \leftarrow \text{priority}(T_1)$
  - $T_2$  blocked by  $T_3$ :  $\text{priority}(T_3) \leftarrow \text{priority}(T_1)$



# Chained Blocking: Problem of PIP



In the worst case, the highest priority task  $T_1$  can be blocked by  $N$  lower priority tasks in the system when  $T_1$  has to access  $N$  semaphores to finish the execution!



# Priority Ceiling Protocol (PCP)

---

- Avoid chained blocking
  - Guarantee a task is blocked by **at most one** lower priority task
  - **No deadlock**
- Assumptions
  - Each task is associated with a fixed priority

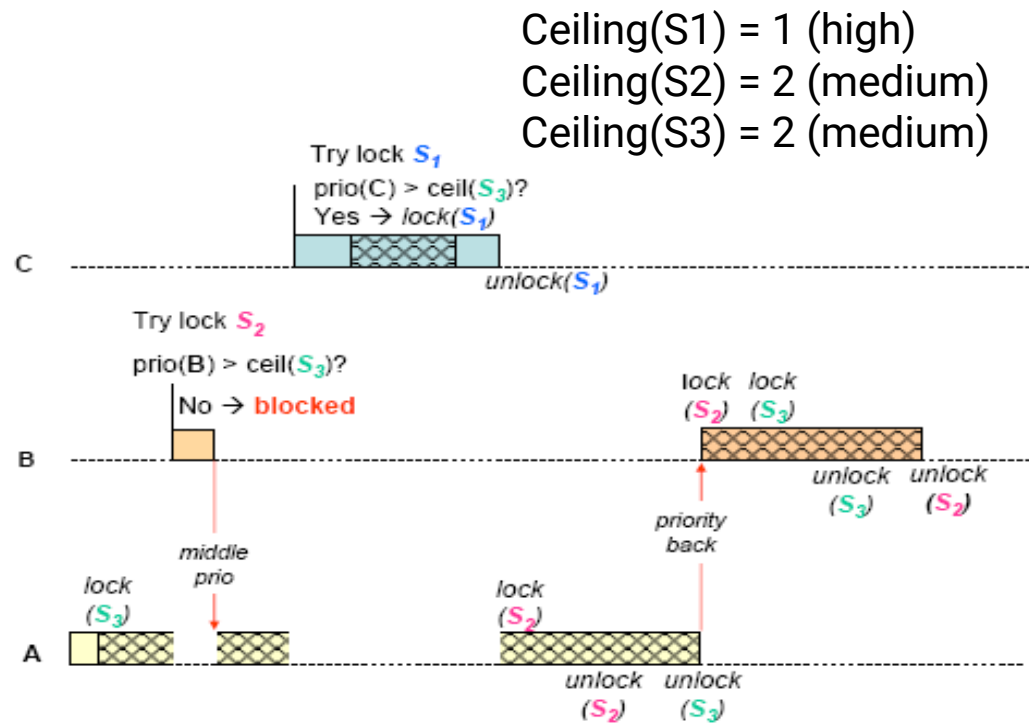


# PCP

---

- Each semaphore has a fixed priority ceiling
  - $\text{ceiling}(S)$  = highest priority among all the tasks that will request  $S$
- How it works:
  - $T_i$  can access a semaphore  $S$  if
    - $S$  is not already allocated to any other task; and
    - Priority of  $T_i$  is higher than the current processor ceiling =  $\max(\text{priority ceilings of all the semaphores allocated to tasks other than } T_i)$

# PCP



No chained blocking!