

CAS CS 552 Intro to Operating Systems

Richard West

Virtual Memory


Objectives

- To describe the benefits of a virtual memory system
- To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames
- To discuss the principle of the working-set model


Background

- **Virtual memory** – separation of user's (logical) view of memory from physical memory
 - Only part of a program needs to be in physical/main memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by multiple processes
 - Allows for more efficient process creation
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation


Virtual Memory Pages Mapped to Physical (Page) Frames



Virtual-address Space



Shared Library Using Virtual Memory




Demand Paging

- Bring a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - Higher degree of multiprogramming
- Page is needed \Rightarrow a reference is made to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
- Lazy swapping** – never swaps a page into memory unless page will be needed
 - Swapper that deals with pages is a **pager**

Page Validation

- Each page table entry has a valid/invalid bit
 - Used to indicate whether or not a page is in memory
 - Let **v** \Rightarrow in-memory, **i** \Rightarrow not-in-memory
 - NOTE: Some architectures refer to the valid bit as a “present” bit
 - Initially valid–invalid bit is set to **i** on all entries

Example Page Table w/ Some Unmapped Pages



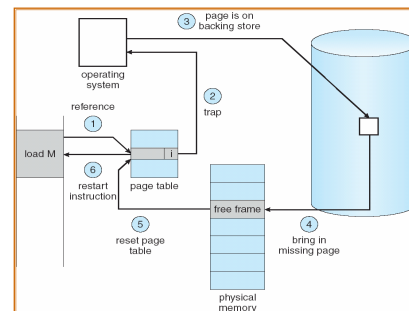
Page Faults

- If a process executes and accesses all pages in memory (i.e., all memory references are to valid pages), execution proceeds normally
- If a process tries to access a page, **P**, not in memory:
 - Paging hardware will notice invalid bit for **P** when translating an address via the active page table
 - This causes a **page fault** trap to the OS

Handling Page Faults

- OS checks info in the process control block (PCB) for the current process causing the page fault, to determine whether or not the memory reference was within the process's address space
 - PCB has info about the contiguous virtual memory areas making up the process's address space
 - If reference was to an address outside process's address space, the process is terminated (a memory protection fault)
 - If reference was within process's address space but page is *not* in memory, page is mapped to a free frame
- The page table is updated with a valid entry for the page that faulted
 - Info about the contiguous VM areas of the process are updated as necessary
- The instruction that caused the page fault is restarted as though it had always been in memory

Steps in Handling a Page Fault



Performance of Demand Paging

- Let $p = \text{Pr}\{\text{page fault}\} \mid 0 \leq p \leq 1.0$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- Effective Access Time (EAT)
$$\text{EAT} = (1 - p) \times \text{memory access time to swap page out, time to swap page in, \& time to restart instruction})$$

Demand Paging Example

- Memory access time = 200 nanoseconds (nS)
- Average page-fault service time = 8 milliseconds
- $$\begin{aligned}\text{EAT} &= (1 - p) \times 200 + p (8 \text{ milliseconds}) \\ &= (1 - p) \times 200 + p \times 8000000 \text{ nS} \\ &= 200 + p \times 7999800 \text{ nS}\end{aligned}$$
- If one access out of 1000 causes a page fault, then
 $\text{EAT} = 8.2 \text{ microseconds}$
This is a slowdown by a factor of about 40
Need to keep page fault rate low

Process Creation

- Virtual memory allows other benefits during process creation:
 - Copy-on-Write
 - Memory-Mapped Files (later)

Copy-on-Write

- Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory

If either process modifies a shared page, only then is the page copied
- To do this, pages have a read-only bit set that causes a trap to the kernel when either parent or child attempts to write to the corresponding page
- COW allows more efficient process creation as only modified pages are copied

Page Replacement 1/2

- The story so far: page-based memory...
 - eliminates external fragmentation
 - enables virtual memory to be larger than physical memory
 - provides memory protection for different address spaces (at granularity of a page)
- To maintain efficient use of CPU, it is desirable to have a *high degree of multiprogramming*
- Increasing degree of multiprogramming increases the # of processes at least partially loaded in memory
 - Nearly all physical memory may be in use
 - If a process requires a page to be loaded into memory (e.g., due to demand paging), there might *not* be any free frames
- In this case, on a page fault, the OS must replace an existing page in memory**


Page Replacement 2/2

- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement
- Use **modified (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk to free up space for new pages

Basic Page Replacement Approach

- Find the location of the desired page on disk
- Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
- Bring the desired page into one of the free frames
- Update the page and frame tables accordingly
- Restart the process


Diagrammatic View of Page Replacement



Page Replacement Algorithms

- Want lowest page-fault rate
- Evaluate algorithm by running it on a particular string of memory references and computing the number of page faults on that string
 - For simplicity, a reference string will refer to a series of page numbers corresponding to the virtual memory addresses being referenced

Typical Page-Fault Distribution



First-In-First-Out (FIFO) Algorithm

- OS maintains a queue (for all currently loaded) pages based on the time they are brought into memory
- Oldest page is at the head of the queue & newest is at the back
 - When a page is replaced, it is taken from the front of the queue
- FIFO replacement can suffer from Belady's Anomaly
 - more frames \Rightarrow more page faults

Example FIFO Page Replacement

- Reference string:
 - 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames:


1	4	5
2	1	3
3	2	4

9 page faults
- 4 frames:


1	5	4	
2	1	5	
3	2		
4	3		

10 page faults

Additional FIFO Page Replacement Example




FIFO Replacement Illustrating Belady's Anomaly



Optimal Algorithm

- Achieves lowest page fault rate
- Aim: replace the page that will not be used for the longest period of time (i.e., whose next reference is furthest in the *future*)
- Problem: Optimal algorithm requires future knowledge of memory, and hence, page references which is not usually possible to know
- If we use the recent past to predict the future, then we can replace the page that has not been used for the longest time
 - Least recently used (LRU) replacement

Optimal Page Replacement




LRU Page Replacement Example 1

- Let there be 4 frames to use
- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

LRU Page Replacement Example 2

- Here, there are 3 frames to use




LRU Implementation 1/2

- (1) Counter implementation:
 - A logical clock (or counter) is incremented every memory reference
 - Each page has a "time of use" field
 - The "time of use" field is updated with the counter value when the page is referenced
 - The page to replace is the one with the lowest counter value
 - Problems:
 - counter overflow
 - must search the counter values for every page to determine lowest counter value

LRU Implementation 2/2

- (2) Stack implementation:
 - keep a stack of page numbers for all loaded pages
 - Most recently referenced page is removed from the stack and placed on top of the same stack
 - Over time, the page at the bottom of the stack is the LRU page
 - Problems: rearranging the stack can require more overhead than the counter to keep track of page references
 - Both stack- and counter-based LRU implementations require h/w assistance in practice, since memory references can be in the multiple millions per second

Example Stack Implementation



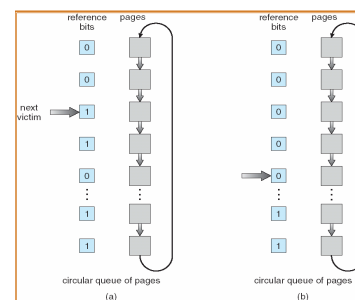
LRU Approximation

- With limited hardware support many computer systems implement alternatives, or approximations, to LRU
 - Some systems use a **reference bit** (for each page) set by the hardware when the page is referenced
 - Can replace a page whose reference bit is not set (if one exists)
 - ...but this loses historical ordering info about page references

Second Chance (Clock) Algorithm

- Like FIFO replacement but first check the reference bit of a page
 - If bit=0, replace page
 - If bit=1, clear bit and check next page
 - This gives page a "second chance" to stay in memory
 - Worst-case: all bits are 1 for all pages in memory
 - Here, we cycle through all pages, clearing their reference bits, and return to the first page in the cycle (which is replaced)

Second-Chance (Clock) Page Replacement Algorithm



Enhanced Second Chance Algorithm

- a.k.a "not recently used", or NRU
- This uses a **reference bit** and a **modify bit** per page
- Each page can be in one of four classes:

	referenced	modified	
0	0	0	Not recently ref'd or modified
0	0	1	
1	1	0	Recently ref'd but not recently modified
1	1	1	

Most desirable to replace

A page may be modified but not recently ref'd if it is awaiting write-back to disk

Counter Algorithms

- Using multiple reference bits (e.g., 8) for each page in memory, we can approximate a finite history of references
 - At regular intervals (e.g., 100mS) a timer interrupt causes the OS to right-shift by 1-bit an n-bit counter for each page, and then insert the reference bit of that page into the most significant bit of the counter
 - Least significant bits are discarded
- e.g. 1, if counter = 00000000 page has not been ref'd for last eight clock intervals
- e.g. 2, a page with counter=10001100 has been more recently ref'd than one page with counter 01111111

Counter Algorithm Variants


- **LFU** – replace page with smallest counter value
 - Problem?
 - A page might be referenced frequently and then not at all but may remain in memory long after it is needed
- **MFU** – replace page w/ largest counter value as smallest one may have only just been brought into memory (and may need to be used again in the future)

Global vs. Local Replacement

- **Global replacement**
 - System may replace a page from any process/address space
- **Local replacement**
 - System replaces a page associated with the local process/address space

Example - Linux


- Paging policy is a variant of the clock algorithm
- Linux maintains **active** and **inactive** lists for all pages in memory
- Pages move through the following states depending on the rate of them being referenced, while at given time intervals pages are aged (i.e., dropped to lower priority states)



Thrashing 1/2

- Suppose a process address space contains two pages, P1 and P2
- Let P2 be brought into memory at the cost of P1 being replaced
 - What happens if the process requires P1 shortly after it is swapped out?
 - Will end up with another page fault
 - If there is a lot of paging activity, a process may spend more time involved in paging than executing
 - Such a process is said to be **thrashing**

Thrashing 2/2



Ways to Limit Thrashing

- Decrease the degree of multiprogramming
- Use only local page replacement algorithms
 - i.e., replace pages from the local process and not other processes
- Provide a process with as many frames as it needs
 - How do we do this?

Locality of Reference

- "90:10" Rule
 - A program spends 90% of its time executing 10% of its code
 - Loops, subroutines, procedures/functions, basic code blocks define "localities of reference"
- If system allocates enough frames to a process for its current locality, page faults will be reduced until the process changes locality


Working Set Model 1/2

- The set of pages in the most recent Δ page references is the **working set**
- The working set is an approximation of the program's locality
 - It is the set of pages currently being used by a process and the entire working set should be in memory to reduce page faults
- Δ should be large enough to encompass a process's current locality, but no larger

Working Set Model 2/2


- Let $\alpha_i(t)$ be the working set size of process i at time, t
- Total demand for page frames from n processes is:
 - $D = \sum_{i=1}^n \alpha_i(t) \mid 1 \leq i \leq n$
- If $D >$ total frames in system, thrashing will occur
- Solution:
 - OS monitors working set for each process and allocates enough frames for the working set corresponding to the current locality
 - If there are spare frames, a new process can begin (thereby increasing the degree of multiprogramming)
 - If D exceeds available frames, OS suspends one or more processes, by writing pages of suspended processes to disk

Working-set model



Page-Fault Frequency Scheme

- To reduce thrashing the OS may take the following approach:
 - Upper and lower bounds can be placed on the desired page-fault rate of a process
 - If a process's page fault rate exceeds the upper bound, it is allocated another frame
 - If a process's page fault rate drops below the lower bound, it loses a frame



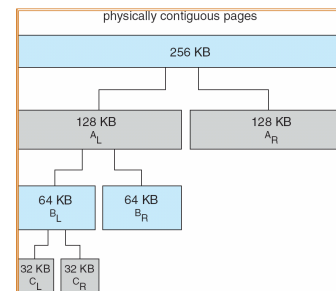
Allocating Kernel Memory

- Treated differently from user memory
- Often allocated from a free-memory pool
 - Kernel requests memory for structures of varying sizes
 - e.g., for file objects and process control blocks
 - Some kernel memory needs to be contiguous
 - e.g., for DMA transfers

Buddy System

- Allocates memory from fixed-size segment consisting of physically-contiguous pages
- Memory allocated using **power-of-2 allocator**
 - Satisfies requests in units sized as power of 2
 - Request rounded up to next highest power of 2
 - When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - Continue until appropriate sized chunk available


Buddy System Allocator



Slab Allocator

- Alternate strategy
- Slab** is one or more physically contiguous pages
- Cache** consists of one or more slabs
- Single cache for each unique kernel data structure
 - Each cache filled with **objects** – instantiations of the data structure
- When cache created, filled with objects marked as **free**
- When structures stored, objects marked as **used**
- If slab is full of used objects, next object allocated from empty slab
 - If no empty slabs, new slab allocated
- Benefits include no fragmentation, fast memory request satisfaction

Slab Allocation





Other Issues -- Prepaging

- Prepaging
 - To reduce the large number of page faults that occurs at process startup
 - Prepage all or some of the pages a process will need, before they are referenced
 - But if prepagged pages are unused, I/O and memory was wasted
 - Assume s pages are prepagged and α (fraction) of the pages are used
 - Is cost of $s * \alpha$ saved pages faults > or < than the cost of prepaging $s * (1 - \alpha)$ unnecessary pages?
 - α near zero \Rightarrow prepaging less beneficial



Other Issues – Page Size

- Page size selection must take into consideration:
 - fragmentation
 - page table size
 - I/O overhead
 - locality



Other Issues – TLB Reach

- TLB Reach - The amount of memory accessible from the TLB
- TLB Reach = (TLB Size) X (Page Size)
- Ideally, the working set of each process is stored in the TLB
 - Otherwise there is a high degree of page faults
- Increase the Page Size
 - This may lead to an increase in fragmentation as not all applications require a large page size
- Provide Multiple Page Sizes
 - This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation



Other Issues – I/O interlock

- **I/O Interlock** – Pages must sometimes be locked into memory
- Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm