

Zago I.M.
U88387934

1.

I tried to figure out this question.

For each task, we have a corresponding weight.

The instantaneous share $f_i(t)$ of process i at time t is
 $f_i(t) = \frac{1}{n}$ since each task have equal weights.

The lag means the difference between the service time a process is entitled to should receive and the actual service time it receives.

Ziqi Tan
U88387934

2.(a) Logical Address:

Selector : offset

$<0x\text{BB88}>$: $<0x\text{3F0000}>$

↓

1011 1011 1000 1000

↓

Index 1011 1011 1000

↓

Out of Limit

Segment Num: 0x1771 \Rightarrow Base Address: 0x00010000 Limit: 0x0001FFFF

(b) $<0x\text{BB98}>$: $<0x\text{3FE00D}>$

↓

1011 1011 100 1000

↓

Index 0x1773 \Rightarrow Base: 0x00030000 Limit: 0x003FFFFF

Linear Address: 0x0042E00D

Page size: 4KB

$$0x0042E00D \div 4\text{KB} = 0x0000\ 042E \Rightarrow 1070$$

(C) $0x003F\ CFA0 \div 4\text{KB} = 0x0000\ 03FC \Rightarrow 1020$

Frame number: 212

Assume frame number starts from 0

$$(212 - 1) \times 4096 - 1 \approx 212 \times 4096$$

$$\approx 212 \times 4096 \approx (212 + 1) \times 4096 - 1$$

$$212 \times 4096 + 0xFAD = 872352$$

Ziqi Tan

U88387934

2. (d)

guest virtual address



page table 1 : 4 times



page table 2 : 4 times



guest physical address : host virtual address

VM thought the page
table as physical, but
in fact it is not.

Answer:

$4+4+4=12$ physical memory
~~references~~

page table 1



page table 2

page table 3

physical address

} 4 physical
memory
~~references~~

If we consider GDT here,
it would be $5+5+5+5=20$ references.

Ziqi Tan
U88387934

3.(a) semaphore lanterns-east = 9
semaphore lanterns-west = 0
semaphore lanternkey-east = 9
semaphore lanternkey-west = 0
mutex lock

procedure east_visitor() {

 wait(lanterns-east);

 wait(key-east); } if (key-east == 0) { signal(key-east); }

// walking-through-house

 signal(lanterns-west);

 signal(key-west);

 signal(chay-west);

}

If a visitor holds a key,
and there is not enough
key for him/her to enter,
then release the resource
he/she holds, which can
avoid visitors on one side
to wait each other's keys.

procedure west_visitor() {

 wait(lanterns-west);

 wait(key-west); } if (key-west == 0) { signal(key-west); }

// walking-through-house

 signal(lanterns-east);

 signal(key-east);

 signal(key-east);

}

(b) In this solution, any number of visitors on any sides will not
lead to a deadlock.

Zigi Tan

U88387934

3. (c)

`pthread_cond_wait (&b->occupied_slot, &b->mutex)`
and

`pthread_mutex_unlock (&b->mutex)` in the check_id

may cause deadlock.

Consider this scenario.

① The buffer is almost full.

② Consumers cannot find their own items.

Consumers unlock the mutex.

③ Producers continue to produce.

Now the buffer is full.

④ The occupied-slot is signaled to wake up one of the consumers.

This consumer still cannot find its own item and then return.

⑤ Now the producer cannot produce item and the call the consumers are waiting for the occupied-slot to be signaled.

A deadlock occurs.

The solution is to let the consumers wake up each other when they cannot find their own items until the corresponding consumer gets its item.

Add "`pthread_cond_signal(occupied_slot)`" right after the "`pthread_mutex_unlock(&b->mutex)`" and right before the "`return NULL`".

Ziqi Tan

U88387934

4. It does not use LOCK to exclusively use the bus in a multicore system.

Fix: Use Add LOCK prefix before bts and btr.

Improve: prevent a CPU who ~~will~~ attempts to acquire the lock continuously use the bus exclusively, which will cause a performance problem on other ^{CPU} cores.

spinlock-lock:

LOCK bts \$0, (mlock)

jc spin-wait

ret

spin-wait:

test (mlock), 1 # if the lock is free

jnz spin-wait # not free

jmp spinlock-lock # may be free

spinlock-unlock:

LOCK btr \$0, (mlock)

ret

Zigi Tar

U88387934

5. (A) 31 16 | 15 0

Base 0:15 | Limit 0:15

b3	5b	55 52	51	48	47	40	39	32
Base	Flags	Limit	Access Byte	Base	16:23			
24:31		16:19						

code segment descriptor: 0x0000FFFF data descriptor: 0x0000FFFF
0x00CFFA00 Access byte: ~~0x00CFF200~~
0x00CFF200

Access byte:

P_r: present bit must be 1.

Privilege: 2 bits. Ring 3: 11

S: Descriptor type. 1 for code or data, 0 for system segments

Executable bit: 1 for ~~code~~ segment and 0 for data.

Direction/Conforming bit:



→ code selector

data selector

1: can be executed from
an equal or lower
privilege level.

0: segment grows up → limit > offset

0: can only be executed
on the ring set specified
above.

1: segment grows down → limit < offset

RW: Readable for code selectors
Writable for data selectors.

Access bit = 0

Flags: ~~00~~

Gr: 0: byte granularity ; 1: 4 kB, Page size

Sz: 0: 16 bit protected mode ; 1: 32 bit protected mode.

Ziqi Tan

U88387934

5. (b) Idt Entry : 0x 0008 0000
0x FCL1 EE00

Reasoning: 31 16 15 0

Segment Selector Pointer

GDT RPL

63 48 47 46 45 44 43 40 39 ~ 32
Pointer P DPL S GateType 0

Present

0xE : 32 bit int gate

Descriptor privilege level: 0 for int and trap gate

which privilege level the calling descriptor minimum should have.

Pointer : FCL1 0000

[5..3] [2] [10]

Segment Selector : 0008 \Rightarrow Index: #
0x01 GDT: 0 RPL: 00

5. (c) No! Not safe!

~~Kernel can never~~

Kernel cannot trust the user-space code, in case that the user-level code is adverse.

e.g.: user level code call foo

foo:

any privileged instruction

lgdt

ret

Kernel can never trust the user level code who wants to change the GDT. The execution flow should pass control to user space first and then execute the user level code, and then check privilege. User level does not have the privilege to lgdt.

Ziqi Tan

U86387934

5. (d) ISR should use iret to return.

Although we try have tried to protect return address,
iret also pops the IP, CS and flag register, which
have been push when an interrupt occurs.

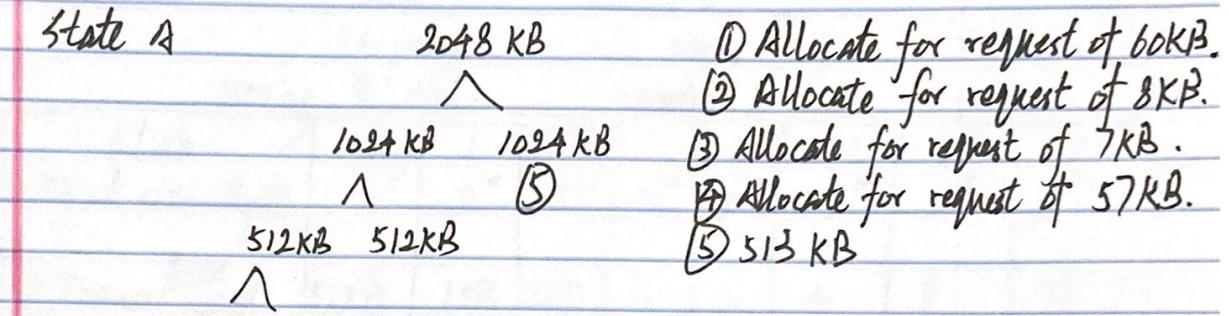
Zizi Tan

U88387934

6.(a) Due to limited space, I use a tree to represent the lists.

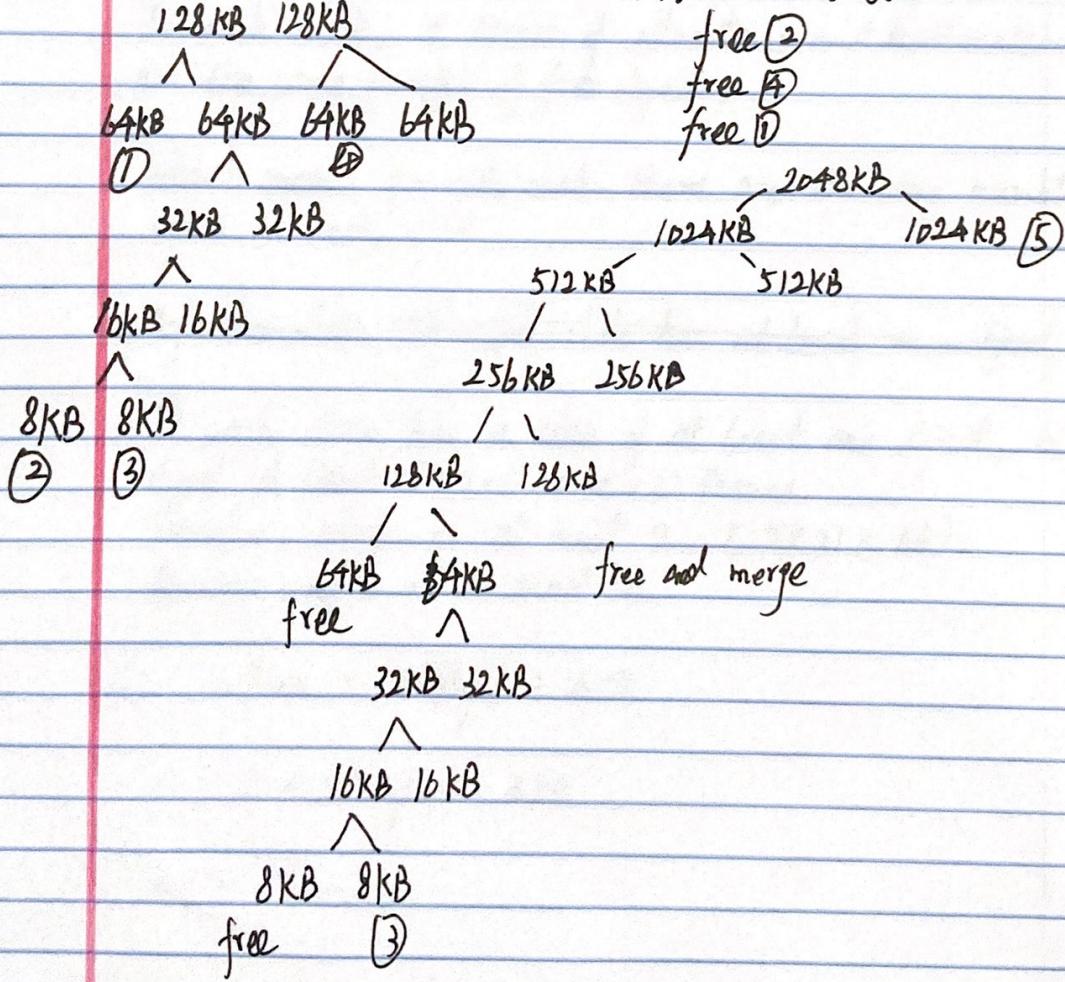
$$2^9 \times 4\text{ KB} = 2048\text{ KB}$$

State A



- (1) Allocate for request of 60KB.
- (2) Allocate for request of 8KB.
- (3) Allocate for request of 7KB.
- (4) Allocate for request of 57KB.
- (5) 513 KB

(b). Deallocations:



Zigi Tam

U88387934

6. (c) Page size : 4 kB
132 Pages

	Cache 1				Cache 2			
	66 pages				66 pages			
	having 8 slabs				having 8 slabs			
Slab	0	1	2	3	4	5	6	7
Object size/Byte	2^9	2^{10}	2^{11}	2^{12}	2^{13}	2^{14}	2^{15}	2^{16}
# Frames/need	512B	1KB	2KB	1	2	4	8	16

i)

In each slab, a series of allocations and deallocations results in the same number of free objects,

which means for each slab, there are the same number of free space in each slab.

~~If the slabs~~ If each slab has at least one object

If each slab has a size of at least one object, slab 3~7 have need at least $1+2+4+8+16=31$ frames.

Therefore, n_{free} is at most 2. ($2 \times 31 \leq 66$)

Answer: n_{free} is at most 2.

ii) $n_{free} \times 2^{16} \text{ Bytes} \div 4 \text{ kB}$

$$= 2 \times 2^{16} \text{ Bytes} \div 4 \text{ kB}$$

$$= 32 \text{ Frames}$$